

RSA® Conference 2019

San Francisco | March 4–8 | Moscone Center



BETTER.

SESSION ID: HTA-T09

The Foreshadow Attack: From a Simple Oversight to a Technological Nightmare

Dr Raoul Strackx, PhD

imec-DistriNet, KU Leuven, Belgium

Foreshadow Attacks

- Independently discovered
- Team of KU Leuven, Belgium
- Team of Universities of Technion, Michigan and Adelaide and DATA61
- Intel discovered other variants

foreshadowattack.eu



Technion
Israel Institute of Technology

UNIVERSITY OF
MICHIGAN



2018 was quite interesting...



source: <https://software.intel.com/security-software-guidance/software-guidance>

**These were vulnerabilities in the processor itself
Hence, virtually every application was effected!**

This led to various reactions...

How we told our upper management at the university (Nov '17)...



Source: <https://pin.it/k4j53t23xiiqcd>

How we told Intel (Jan '18)...



Source: <https://pin.it/k4j53t23xiiqcd>

How IT professionals reacted (to this class of vulnerabilities)...



Source: <https://pin.it/hehzyfhdsvnk>



How do these attacks work, in general?

Side-Channel Attacks



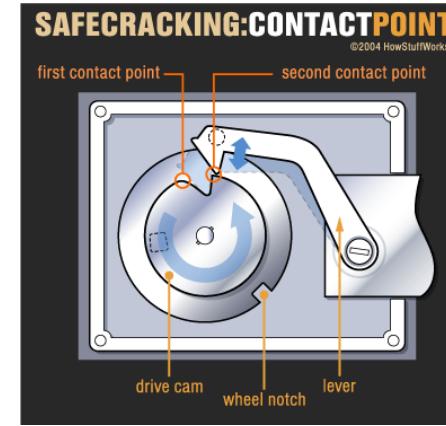
Source: The Italian Job (imdb.com)

Side-Channel Attacks

Attacker



Victim



Action: rotate

Carrier: sound

Charlize Theron

Vault

Security flaw: Lever may produce sound

Sources: <https://home.howstuffworks.com>, [imdb.com](https://www.imdb.com)

How does the Foreshadow attack work?

Attacks

- **Foreshadow-OS**
- **Foreshadow-VMM**
- **Foreshadow-SGX**
- **Foreshadow-SMM**
- →All exploit the same vulnerability, but the attack model varies

Foresight-OS: Reading L1 data through *bare-metal* not-present pages

Side-Channel Attacks

Attacker

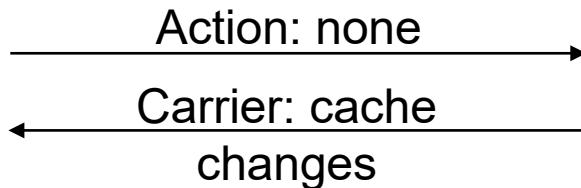


Foreshadow-OS

Victim



Other process'
memory



Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

DistrINet

Side-Channel Attacks

Attacker

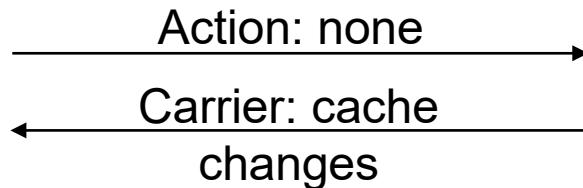


Foreshadow-OS

Victim



Other process'
memory



Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

DistrINet

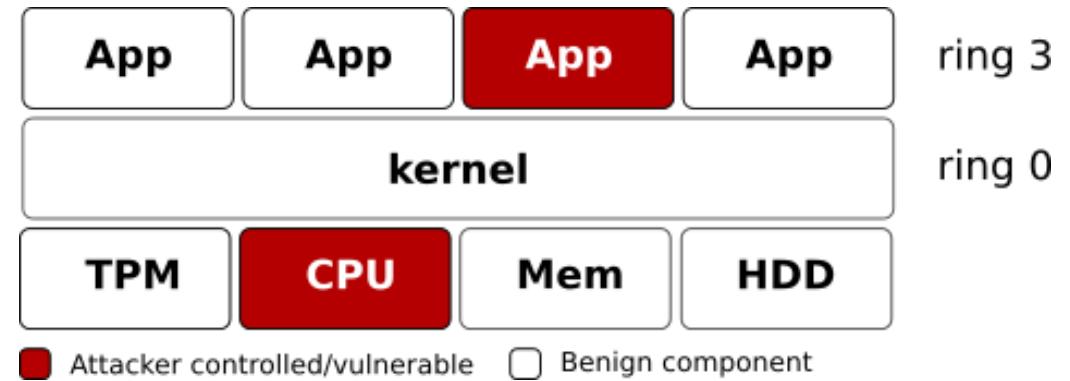
Setting: Attacker-controlled process

Attack model:

- Attacker operates within a malicious process
- Benign, bare-metal kernel ensures process isolation

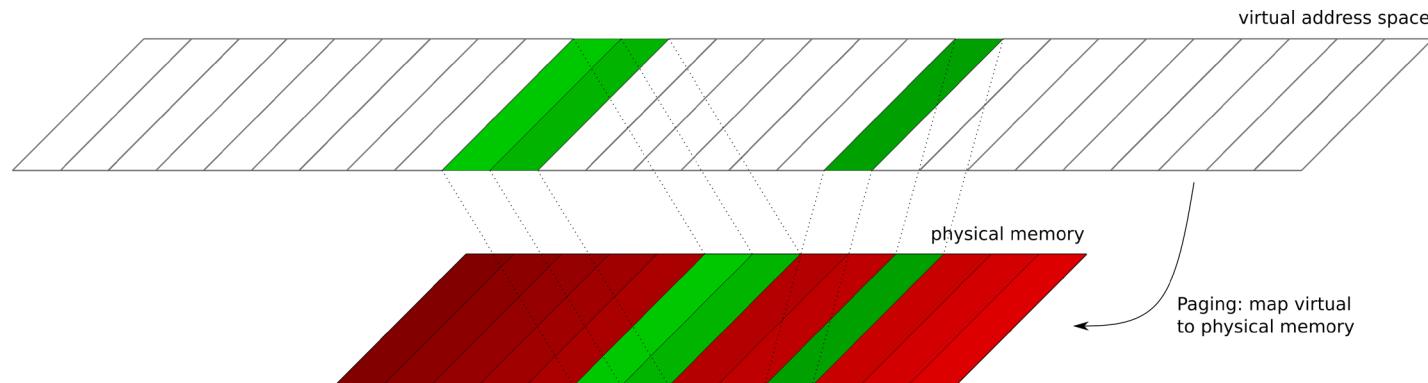
Attack objective:

- Read data outside the process' address space



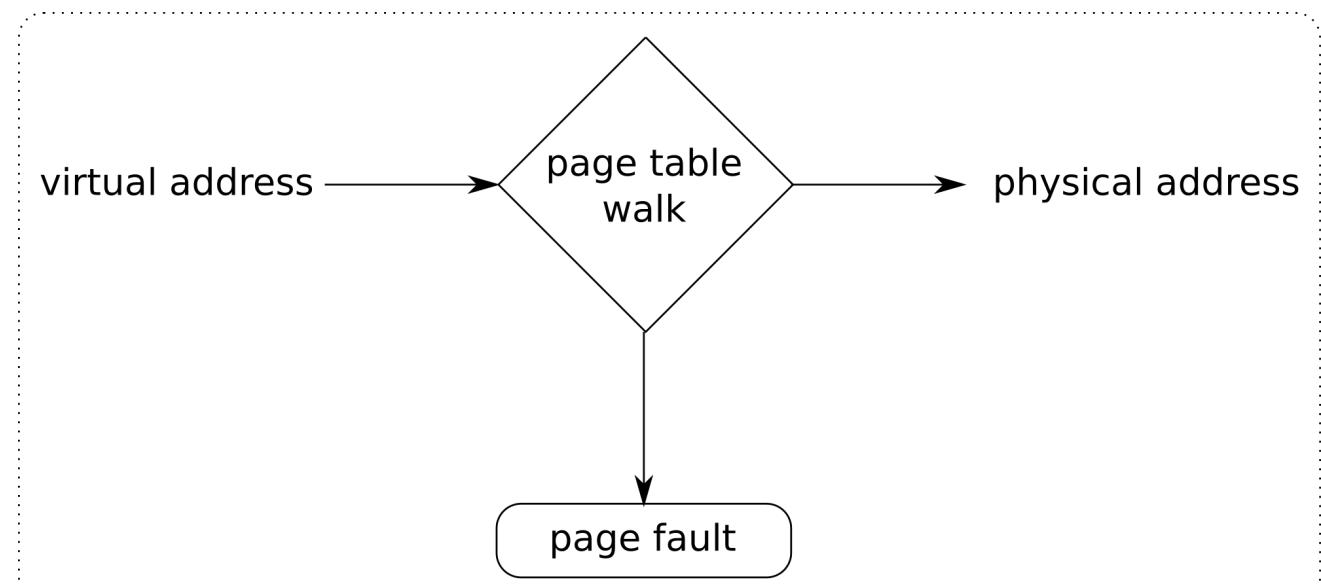
Background: How does process isolation work

- MMU: map virtual address space to physical memory
- Protect physical memory by:
 - Not providing a mapping
 - Restricting access (e.g., U/S-bit)



Background: How does process isolation work

- MMU: map virtual address space to physical memory
- Protect physical memory by:
 - Not providing a mapping
 - Restricting access (e.g., U/S-bit)



Background: How does process isolation work

Bit Position(s)	Contents
0 (P)	Present; must be 1 to map a 4-KByte page
1 (R/W)	Read/write; if 0, writes may not be allowed to the 4-KByte page referenced by this entry (see Section 4.6)
2 (U/S)	User/supervisor; if 0, user-mode accesses are not allowed to the 4-KByte page referenced by this entry (see Section 4.6)
3 (PWT)	Page-level write-through; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
4 (PCD)	Page-level cache disable; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
5 (A)	Accessed; indicates whether software has accessed the 4-KByte page referenced by this entry (see Section 4.8)
6 (D)	Dirty; indicates whether software has written to the 4-KByte page referenced by this entry (see Section 4.8)
7 (PAT)	Indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
8 (G)	Global; if CR4.PGE = 1, determines whether the translation is global (see Section 4.10); ignored otherwise
11:9	Ignored
(M-1):12	Physical address of the 4-KByte page referenced by this entry
51:M	Reserved (must be 0)
58:52	Ignored
62:59	Protection key; if CR4.PKE = 1, determines the protection key of the page (see Section 4.6.2); ignored otherwise
63 (XD)	If IA32_EFER.NXE = 1, execute-disable (if 1, instruction fetches are not allowed from the 4-KByte page controlled by this entry; see Section 4.6); otherwise, reserved (must be 0)

Source: Intel 64 and IA-32 architectures software developer's manual

Background: How does process isolation work

Bit Position(s)	Contents
0 (P)	Present; must be 1 to map a 4-KByte page
1 (R/W)	Read/write; if 0, writes may not be allowed to the 4-KByte page referenced by this entry (see Section 4.6)
2 (U/S)	User/supervisor; if 0, user-mode accesses are not allowed to the 4-KByte page referenced by this entry (see Section 4.6)
3 (PWT)	Page-level write-through; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
4 (PCD)	Page-level cache disable; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
5 (A)	Accessed; indicates whether software has accessed the 4-KByte page referenced by this entry (see Section 4.8)
6 (D)	Dirty; indicates whether software has written to the 4-KByte page referenced by this entry (see Section 4.8)
7 (PAT)	Indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
8 (G)	Global; if CR4.PGE = 1, determines whether the translation is global (see Section 4.10); ignored otherwise
11:9	Ignored
(M-1):12	Physical address of the 4-KByte page referenced by this entry
51:M	Reserved (must be 0)
58:52	Ignored
62:59	Protection key; if CR4.PKE = 1, determines the protection key of the page (see Section 4.6.2); ignored otherwise
63 (XD)	If IA32_EFER.NXE = 1, execute-disable (if 1, instruction fetches are not allowed from the 4-KByte page controlled by this entry; see Section 4.6); otherwise, reserved (must be 0)

Source: Intel 64 and IA-32 architectures software developer's manual

Background: How does process isolation work

"When P-bit is 0, the entry's physical address field may be reused to keep track of the swapped out page"

Bit Position(s)	Contents
0 (P)	Present; must be 1 to map a 4-KByte page
1 (R/W)	Read/write; if 0, writes may not be allowed to the 4-KByte page referenced by this entry (see Section 4.6)
2 (U/S)	User/supervisor; if 0, user-mode accesses are not allowed to the 4-KByte page referenced by this entry (see Section 4.6)
3 (PWT)	Page-level write-through; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
4 (PCD)	Page-level cache disable; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
5 (A)	Accessed; indicates whether software has accessed the 4-KByte page referenced by this entry (see Section 4.8)
6 (D)	Dirty; indicates whether software has written to the 4-KByte page referenced by this entry (see Section 4.8)
7 (PAT)	Indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
8 (G)	Global; if CR4.PGE = 1, determines whether the translation is global (see Section 4.10); ignored otherwise
11:9	Ignored
(M-1):12	Physical address of the 4-KByte page referenced by this entry
51:M	Reserved (must be 0)
58:52	Ignored
62:59	Protection key; if CR4.PKE = 1, determines the protection key of the page (see Section 4.6.2); ignored otherwise
63 (XD)	If IA32_EFER.NXE = 1, execute-disable (if 1, instruction fetches are not allowed from the 4-KByte page controlled by this entry; see Section 4.6); otherwise, reserved (must be 0)

Side-Channel Attacks

Attacker

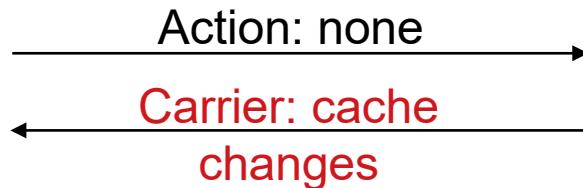


Foreshadow-OS

Victim



Other process'
memory



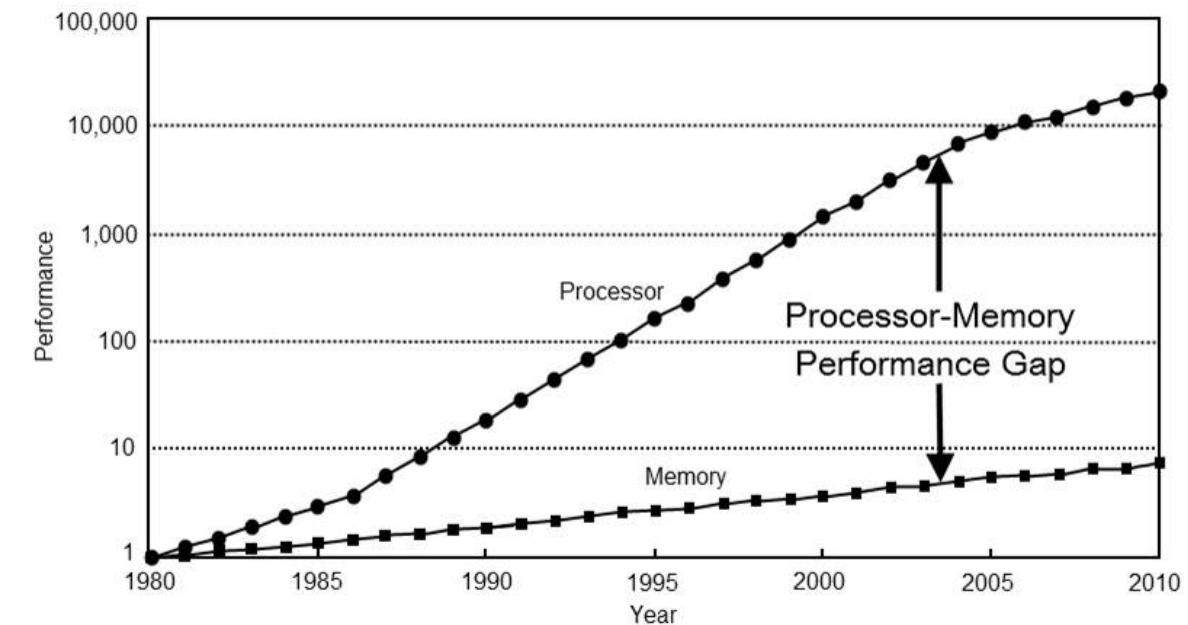
Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

Distrinet

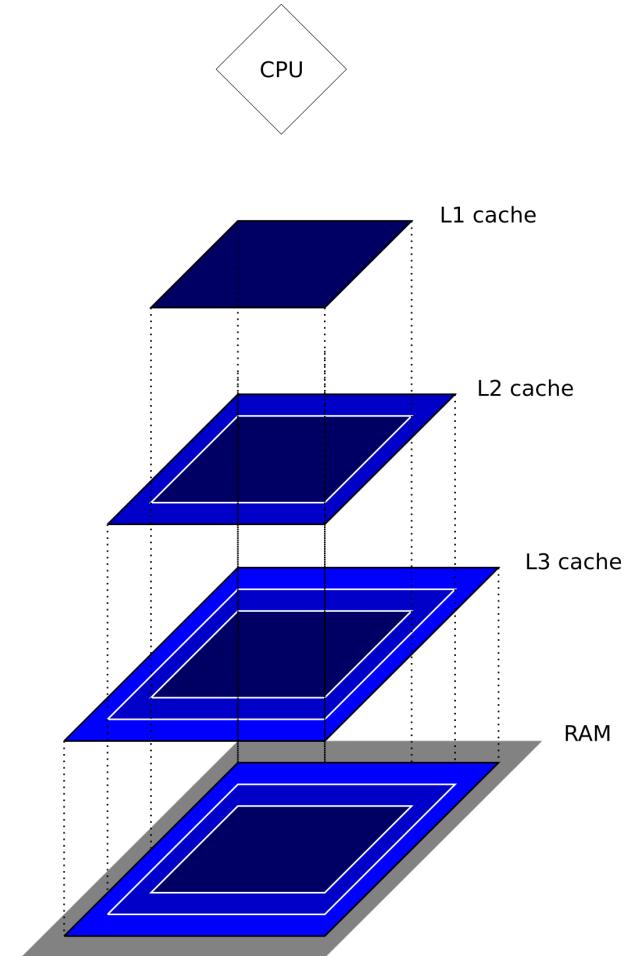
The message carrier: How does the cache work?

- **Problem:** Memory performance grows much slower than CPU performance
- **Solution:** Fast but small caches
 - Intel 486: L1 cache ('89)
 - Pentium Pro: L1 & L2 cache ('95)
 - Today: L1, L2 & L3 caches



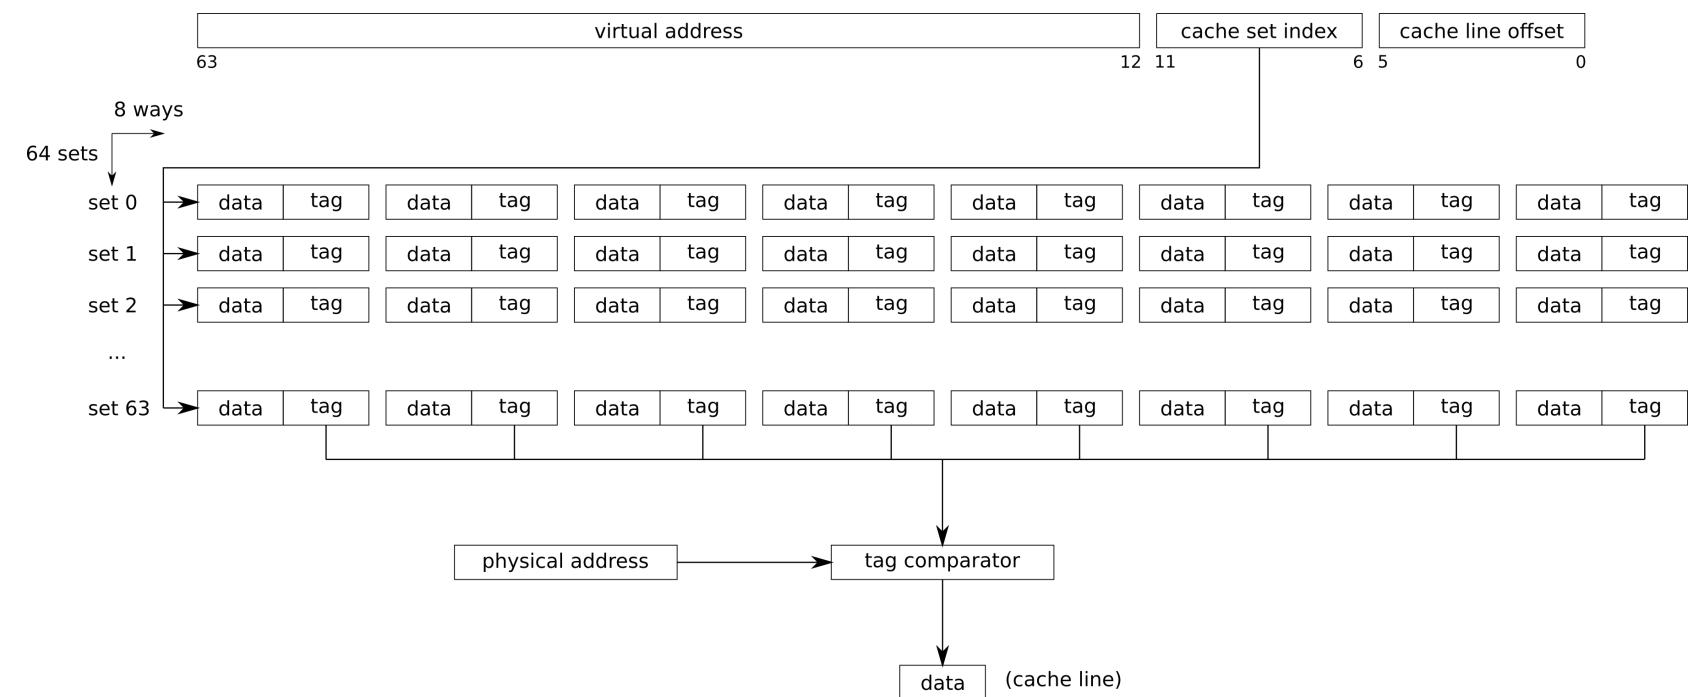
The message carrier: How does the cache work?

- **Problem:** Memory performance grows much slower than CPU performance
- **Solution:** Fast but small caches
 - Intel 486: L1 cache ('89)
 - Pentium Pro: L1 & L2 cache ('95)
 - Today: L1, L2 & L3 caches



The message carrier: How does the cache work?

- Cache lines: 64B
- L1: virtually-indexed, physically tagged
- 64 sets, 8 ways



The message carrier: How does the cache work?

Manipulating the cache:

- Data accesses: load in all cache levels
- *clflush*: Flush data from all cache levels

memory	timing (in cycles)	std. dev.
L1	46	1.25
L2	53	1.14
RAM	246	6.22

→ Any timing results <146 cycles clearly hits the cache

Side-Channel Attacks

Attacker

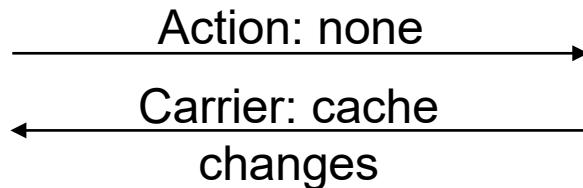


Foreshadow-OS

Victim



Other process'
memory



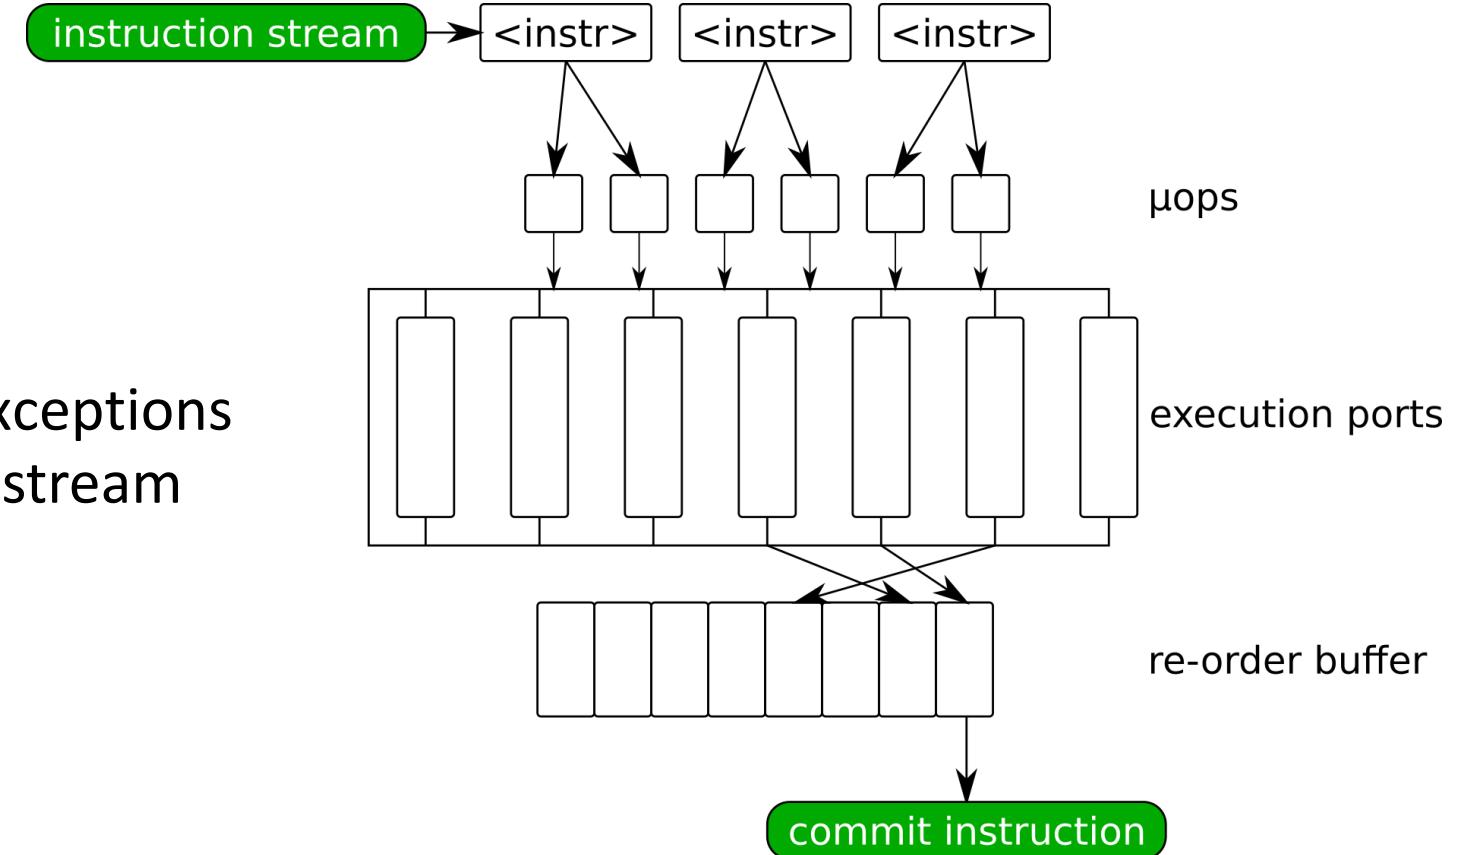
Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

DistrINet

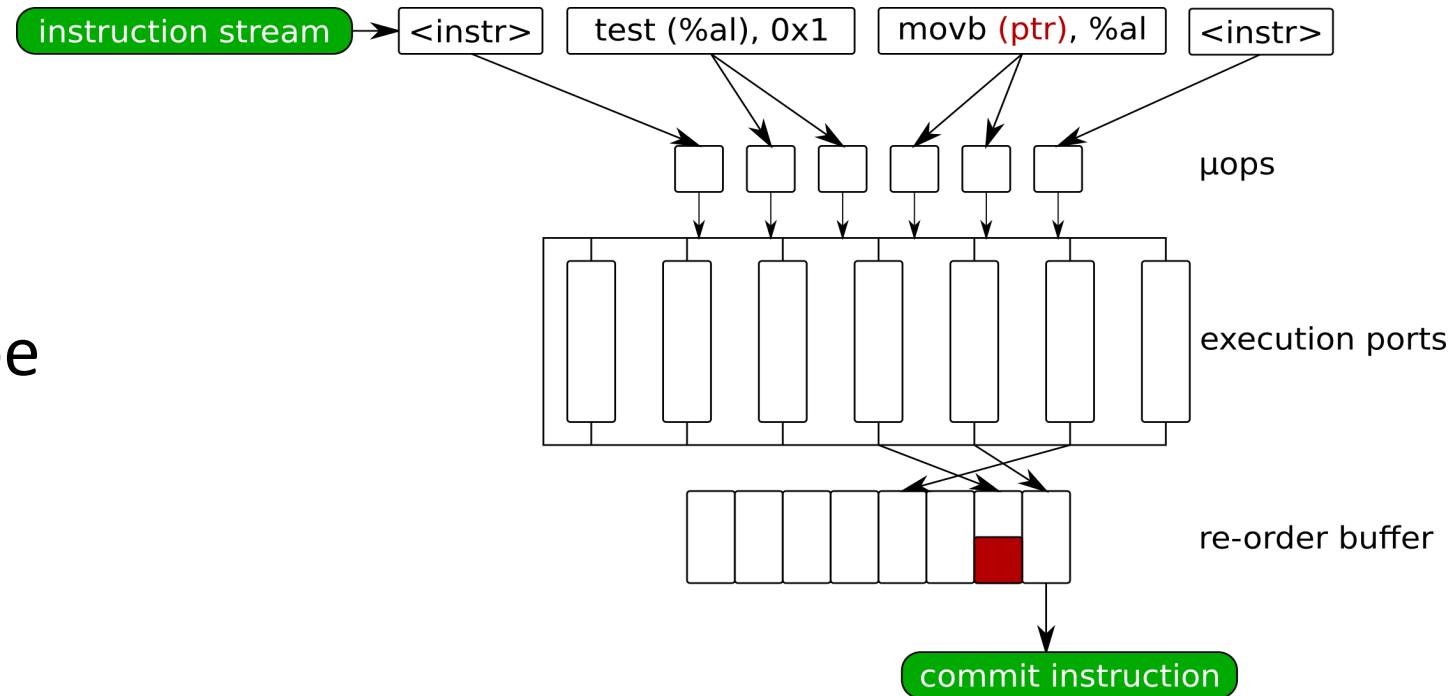
Background: Pipelining & Out of Order Execution

- Split instruction in μ ops
- Use multiple execution ports
- Execute μ ops as soon as possible
- Reorder buffer ensures results/exceptions are visible in-order of instruction stream



The Security Flaw: Transient Execution

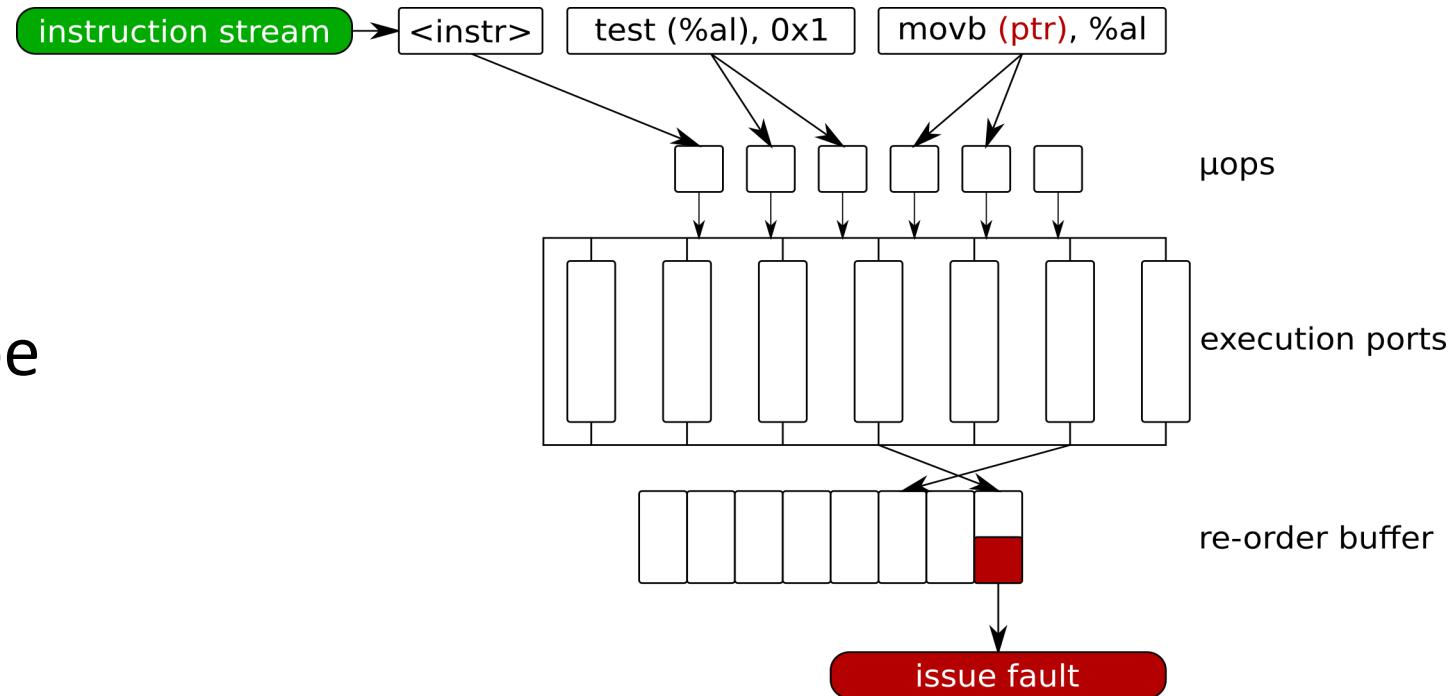
- Transient execution:
- Faults are detected at last moment
- Instruction that should *never* be executed, may already have started



Key Issue: Not all side-effects of transient instructions are rolled back correctly! (e.g., cache changes)

The Security Flaw: Transient Execution

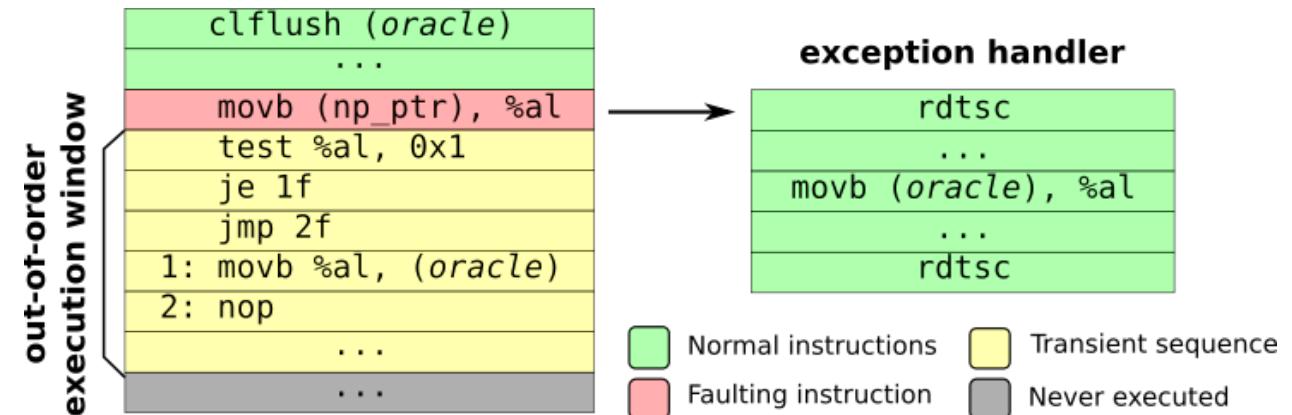
- Transient execution:
- Faults are detected at last moment
- Instruction that should *never* be executed, may already have started



Key Issue: Not all side-effects of transient instructions are rolled back correctly! (e.g., cache changes)

The Security Flaw: Transient Execution

- Transient execution:
- Faults are detected at last moment
- Instruction that should *never* be executed, may already have started



Key Issue: Not all side-effects of transient instructions are rolled back correctly! (e.g., cache changes)

Side-Channel Attacks

Attacker

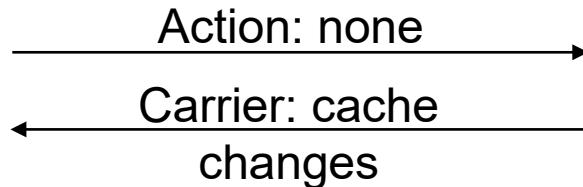


Foreshadow-OS

Victim



Other process'
memory

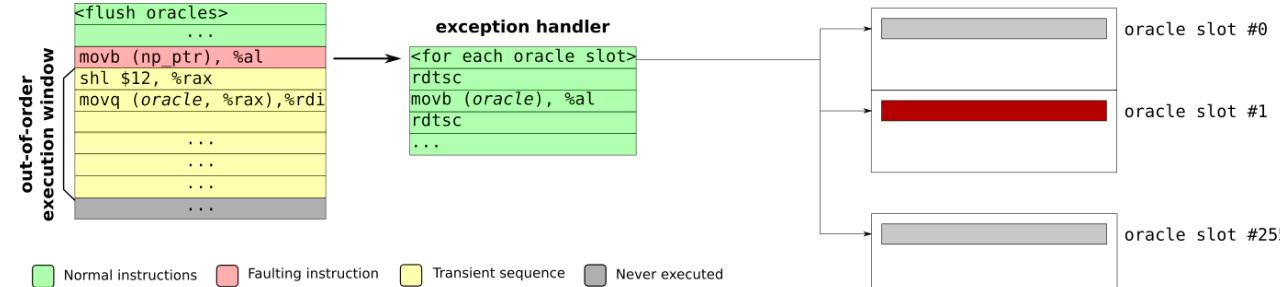


Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

Distrinet

Foreshadow-OS: The exploit

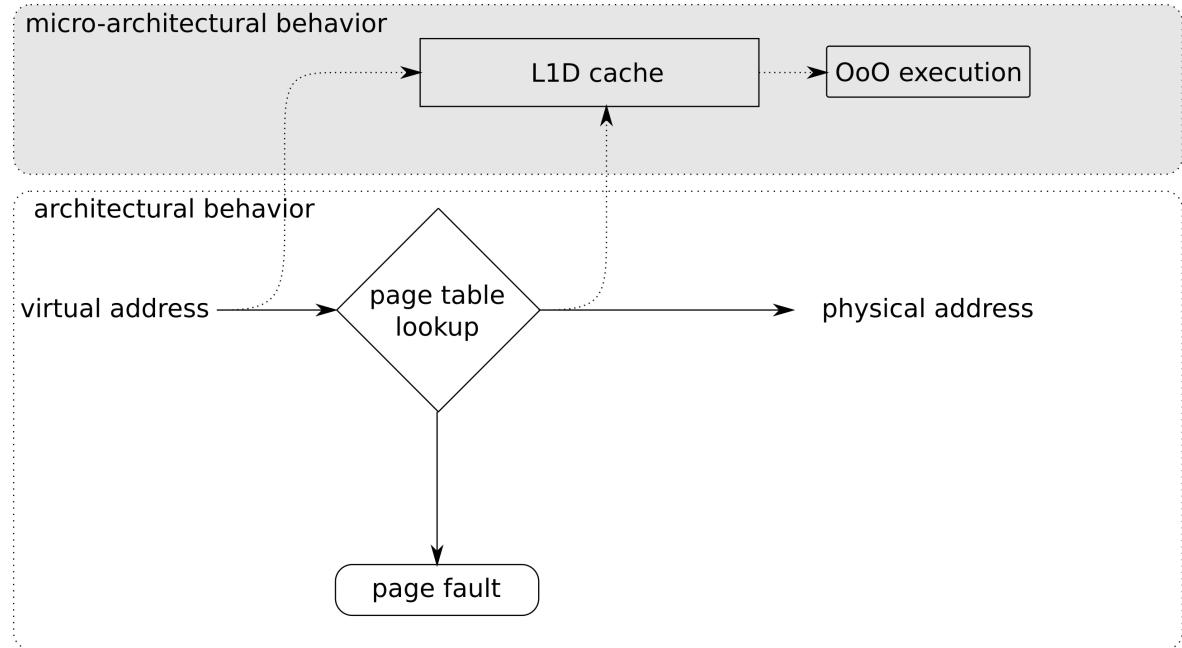


```

1 int8_t *oracles = ...;
2 int8_t *np_ptr = ...; // the secret
3 int8_t _tmp;
4
5 // Step 1: Remove oracle slots from cache
6 for ( int i = 0; i < 256; ++i )
7     clflush( &oracles[4096 * i] );
8
9 // Step 2: Trick system in sensitive data in L1 but PTE present bit to 0
10
11 // Step 3: attempt to read not present memory
12 _tmp = oracle[4096 * (*np_ptr)];
13
14 // suppress fault
15
16 // Step 4: which oracle slot is cached?
17 for ( int i = 0; i < 256; ++i ) {
18     if ( time_access( oracle[4096 * i] ) < 146 )
19         print( "*np_ptr = %i\n", i );
20 }

```

Foreshadow-OS: The exploit



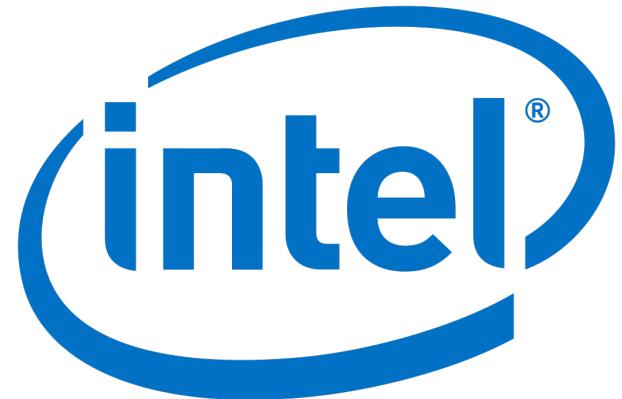
```

1 int8_t *oracles = ...;
2 int8_t *np_ptr = ...; // the secret
3 int8_t _tmp;
4
5 // Step 1: Remove oracle slots from cache
6 for ( int i = 0; i < 256; ++i )
7     clflush( &oracles[4096 * i] );
8
9 // Step 2: Trick system in sensitive data in L1 but PTE present bit to 0
10
11 // Step 3: attempt to read not present memory
12 _tmp = oracle[4096 * (*np_ptr)];
13
14 // suppress fault
15
16 // Step 4: which oracle slot is cached?
17 for ( int i = 0; i < 256; ++i ) {
18     if ( time_access( oracle[4096 * i] ) < 146 )
19         print( "*np_ptr = %i\n", i );
20 }

```

Who's Affected?

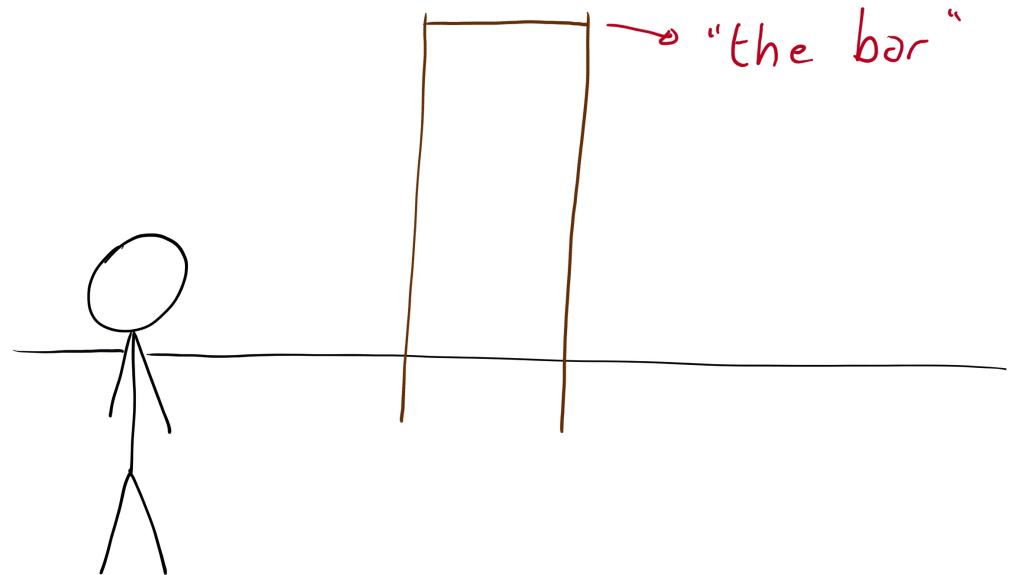
- Intel Core processors of the last 7 years
- Intel server processors
- NOT AMD, not ARM



Impact of This Attack

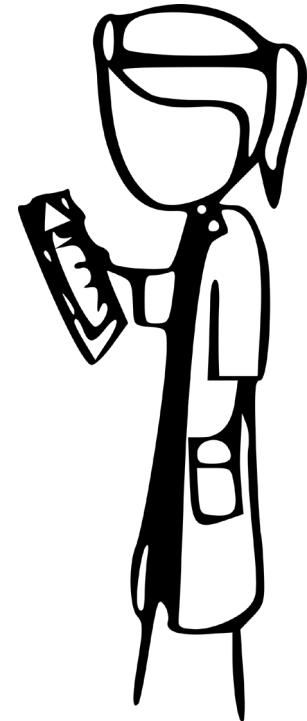
Requirements:

- Secret data in L1D
- Page must be not-present
- ↗ Must difficult attack, “easiest” to understand
- ↗ Low impact!



Mitigations

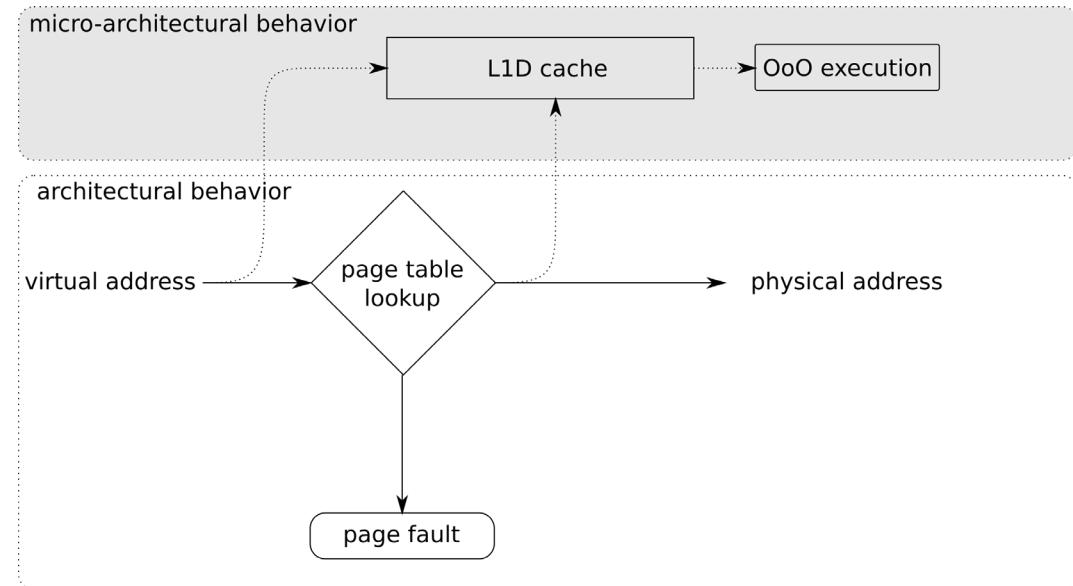
- Long term: Replace chips!
- Short term:
 - Software approaches:
 - Ensure PTE entry do not point to existing physical address
 - Flush complete L1D cache



Source: xkcd.com

Mitigations

- Long term: Replace chips!
- Short term:
 - Software approaches:
 - Ensure PTE entry do not point to existing physical address
 - Flush complete L1D cache



Foreshadow-VMM: Reading physical L1 data through *virtualized not-present pages...*

Side-Channel Attacks

Attacker



Foreshadow-VMM

Victim



Other VM's memory

Action: manipulate PT
Carrier: cache
changes

Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

DistrINet

Side-Channel Attacks

Attacker



Foreshadow-VMM

Victim



Other VM's memory

Action: manipulate PT
Carrier: cache
changes

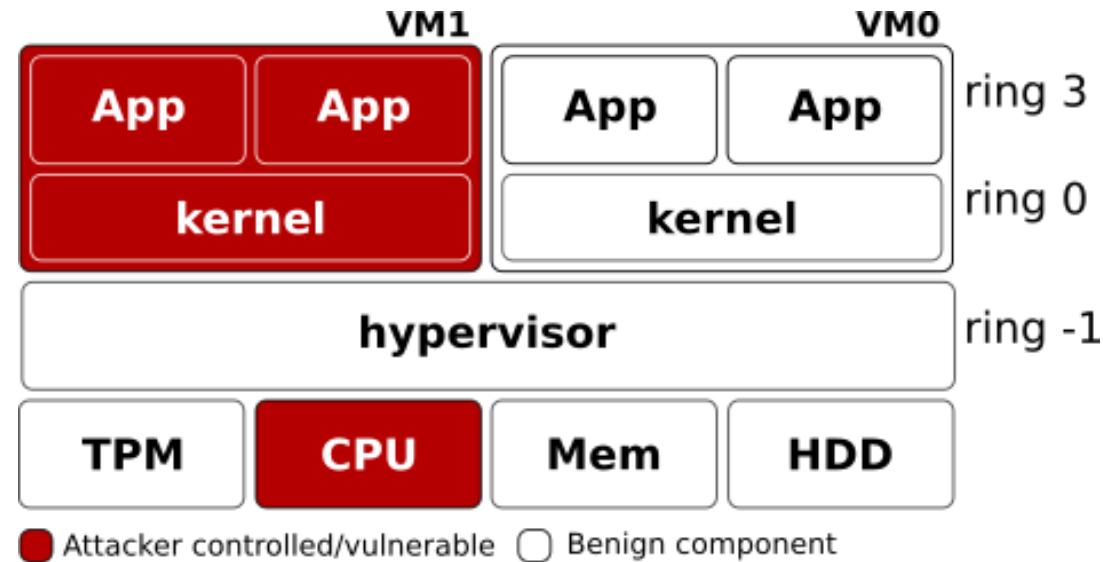
Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

DistrINet

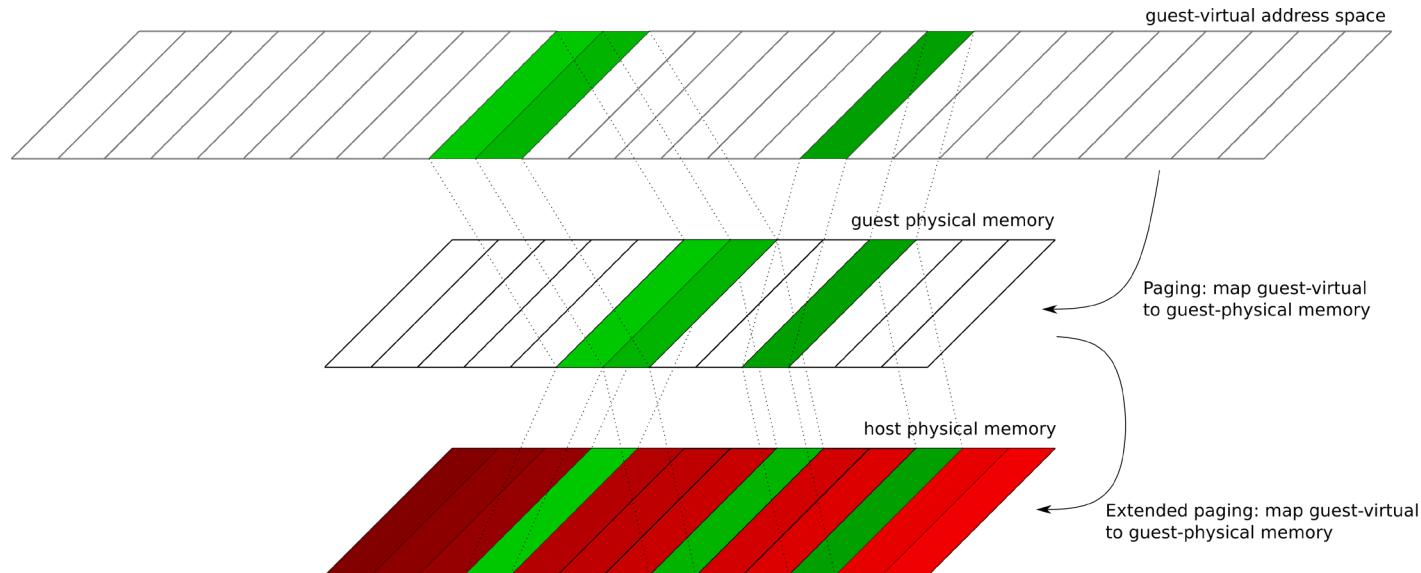
Setting: Attacker-controlled VM

- Multiple VMs on one physical server
- Attacker-controlled VM
- Hypervisor ensures VM isolation
- ↗ Goal: read other VM's memory



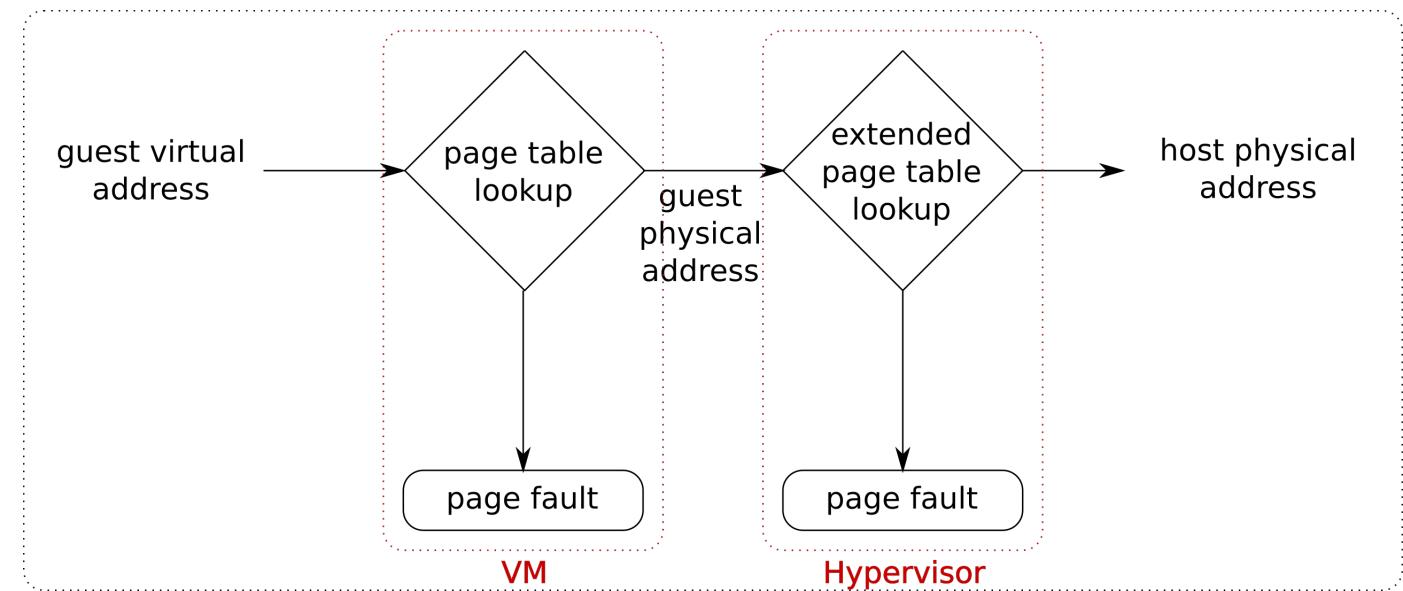
How do extended page tables work

- Adds another layer:
 - PT: guest-virtual addr.
→ guest-phys. addr.
 - EPT: guest-phys. addr.
→ host-phys. addr.



How do extended page tables work

- Adds another layer:
 - PT: guest-virtual addr.
→ guest-phys. addr.
 - EPT: guest-phys. addr.
→ host-phys. addr.



Side-Channel Attacks

Attacker



Foreshadow-VMM

Victim



Other VM's memory

Action: manipulate PT
Carrier: cache
changes

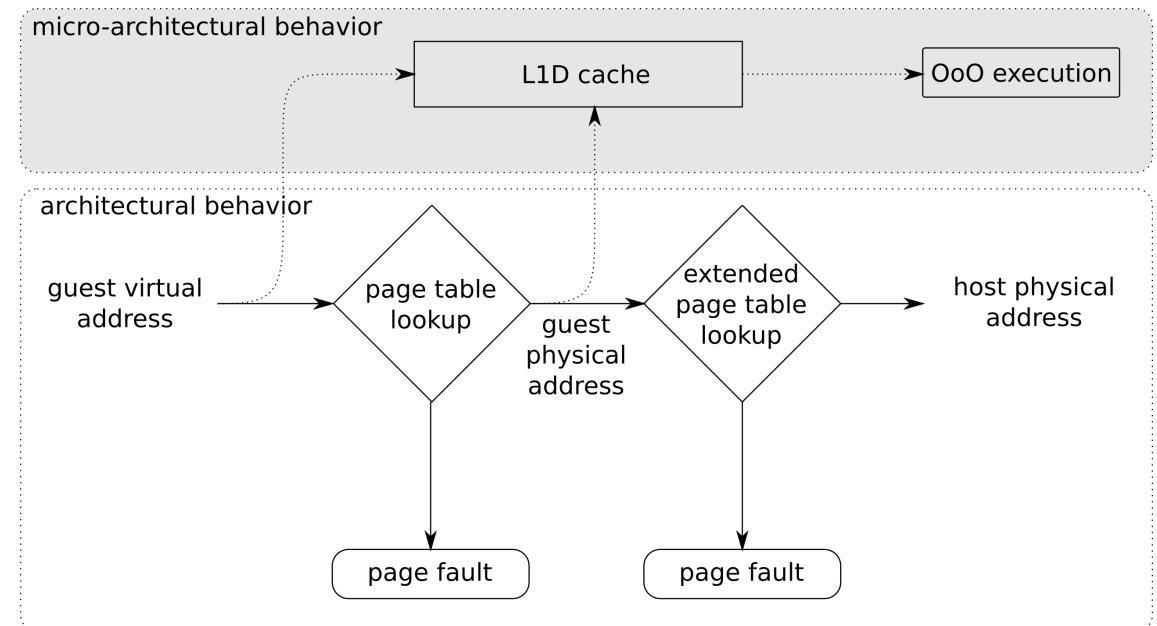
Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

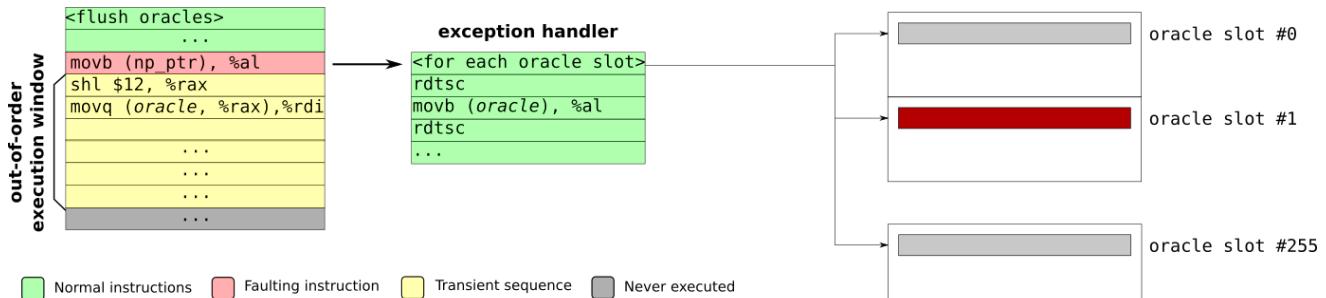
Distrinet

The Security Flaw: Interpreting guest-physical as host-physical addresses

- VM-level: non-present PTE entry
- VMM-level: irrelevant
- Upon access:
 - Tag data access as a violation
 - Pass *guest-physical* as *host-physical* address to L1D cache
 - Continue transient execution!!
- This breaks the VM's address space abstraction!



Foreshadow-VM: The exploit



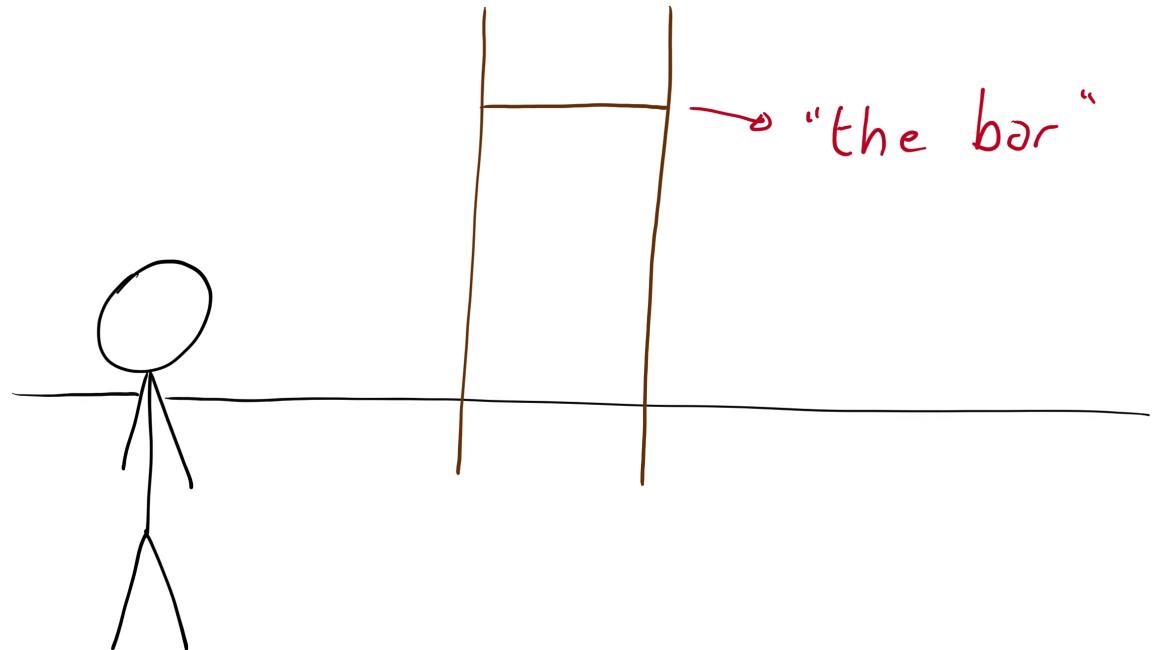
```

1 int8_t *oracles = ...;
2 int8_t *np_ptr = ...;
3 int8_t _tmp;
4
5 // Step 1: Setup PT to physical address of interest
6
7 // Step 2: Remove oracle slots from cache
8 for ( int i = 0; i < 256; ++i )
9     clflush( &oracles[4096 * i] );
10
11 // Step 3: Wait for sensitive data in L1D
12
13 // Step 4: attempt to read not present memory
14 _tmp = oracle[4096 * (*np_ptr)];
15
16 // suppress fault
17
18 // Step 5: is oracle cached?
19 for ( int i = 0; i < 256; ++i ) {
20     if ( time_access( oracle[4096 * i] ) < 146 )
21         print( "*np_ptr = %i\n", i );
22 }

```

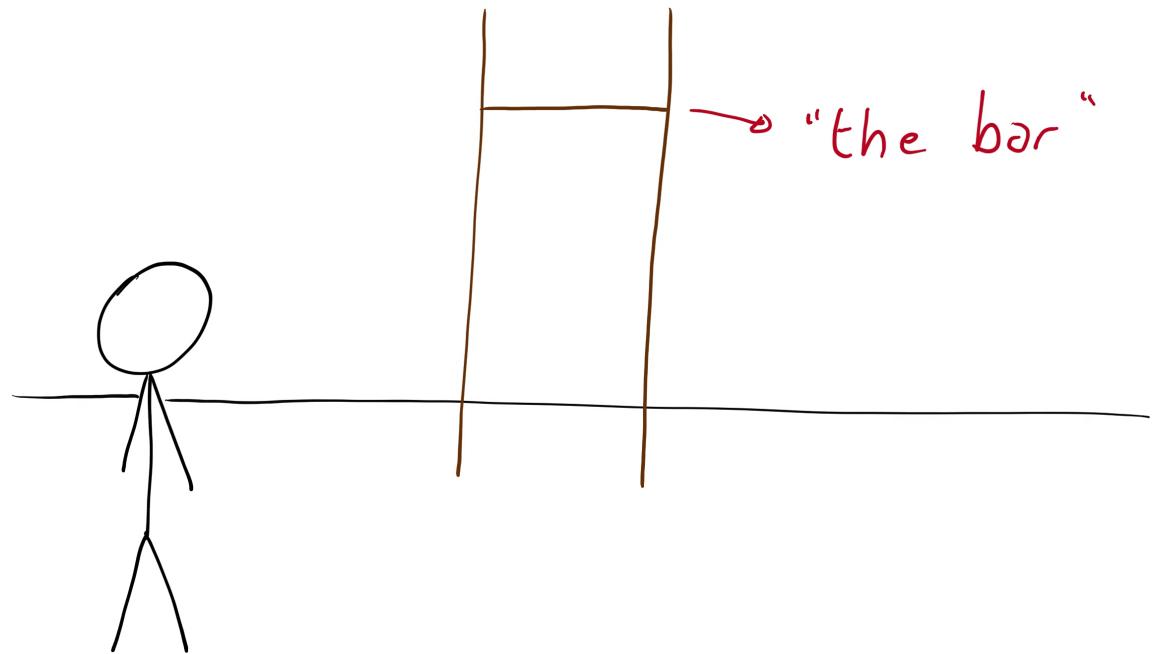
Impact of Foreshadow-VMM

- Requirements:
 - Attacker must have full VM under her control
 - Secret data must reside in L1D
- ↗ Modest impact!



Impact of Foresight-VMM

- Requirements:
 - Attacker must have full VM under her control
 - Secret data must reside in L1D
→ This may not be that complicated!!
- ↪ Modest impact!



Side-Channel Attacks

Attacker



Foreshadow-VMM

Victim



Other VM's memory

Action: manipulate PT
Carrier: cache
changes

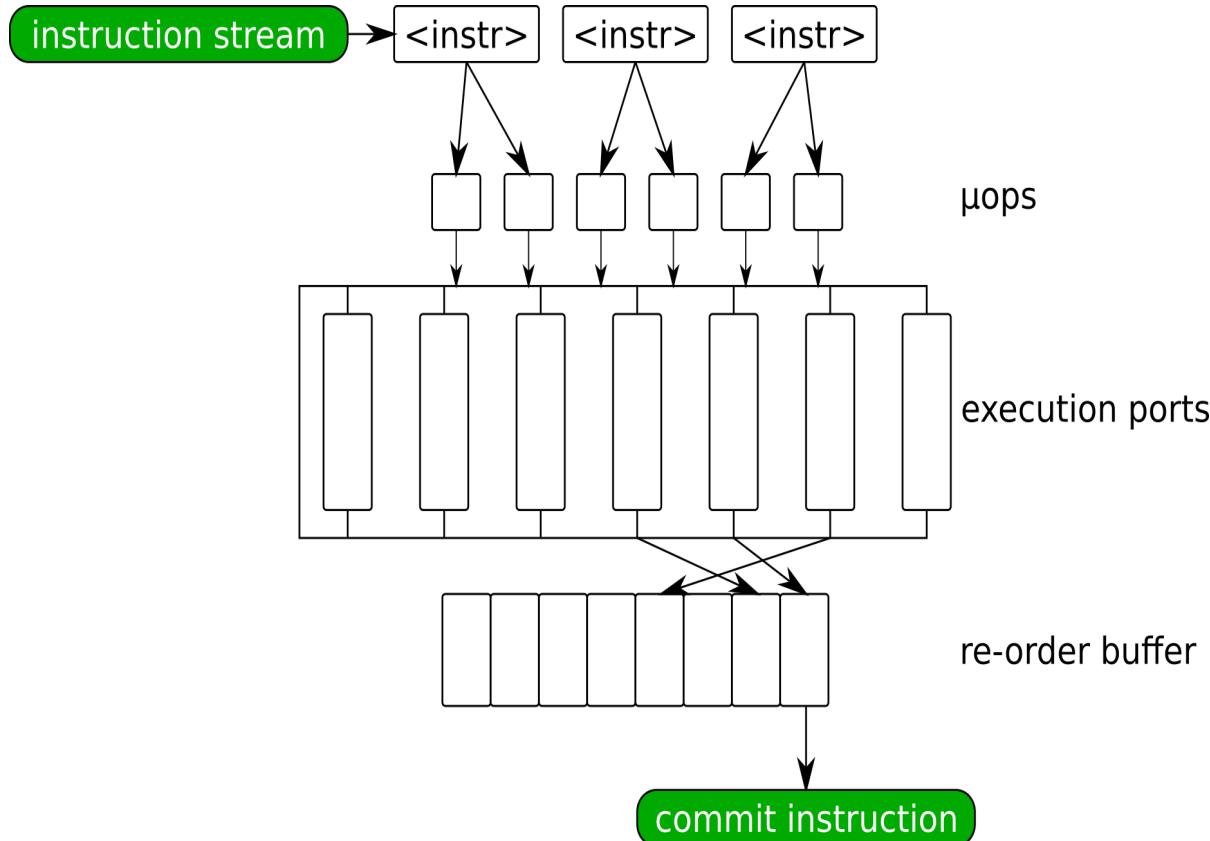
Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

DistrINet

Intel HyperThreading as an Enabler

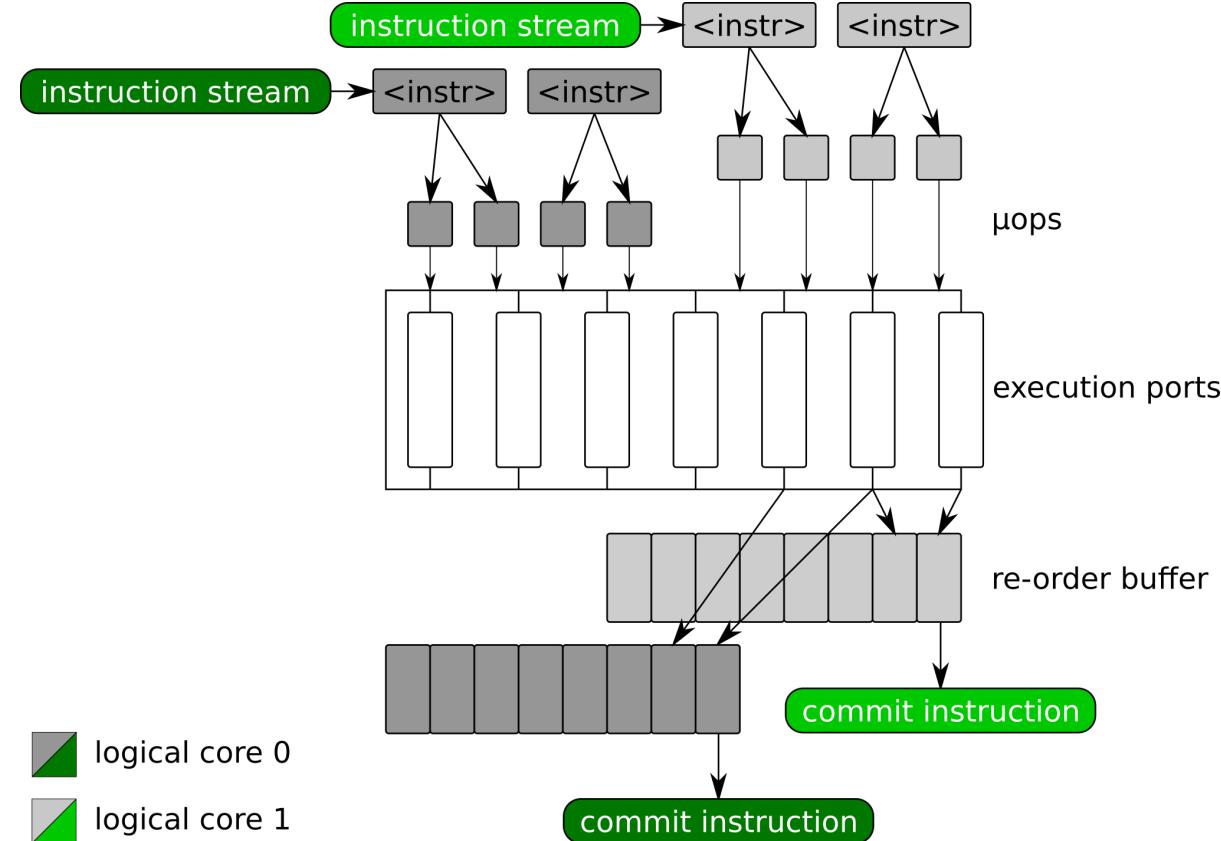
- **Problem:** Execution ports are still under-utilized
- **Solution:** Split physical core in two, duplicate
 - register file
 - re-order buffer
- Shared:
 - Execution ports
 - **L1 cache!** (and other levels)
- ↗ Up to 30% performance increase



source: <https://www.cs.sfu.ca/~fedorova/Teaching/CMPT886/Spring2007/papers/hyper-threading.pdf>

Intel HyperThreading as an Enabler

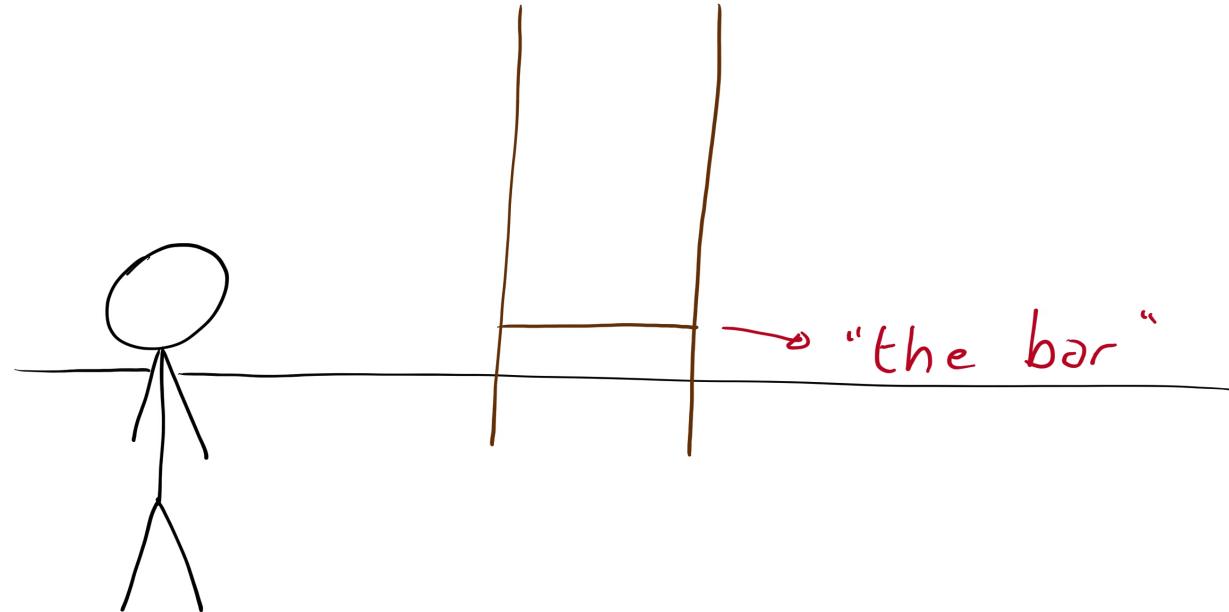
- **Problem:** Execution ports are still under-utilized
- **Solution:** Split physical core in two, duplicate:
 - register file
 - re-order buffer
- Shared:
 - Execution ports
 - **L1 cache!** (and other levels)
- ↗ Up to 30% performance increase



source: <https://www.cs.sfu.ca/~fedorova/Teaching/CMPT886/Spring2007/papers/hyper-threading.pdf>

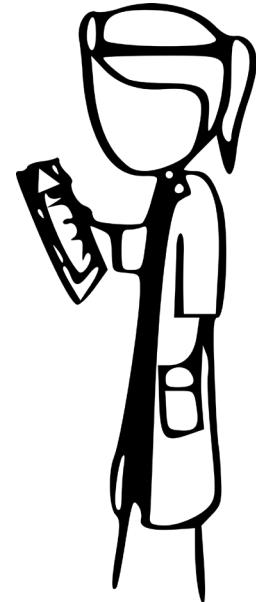
Impact of Foreshadow-VMM

- Requirements:
 - Attacker must have full VM under her control
 - Secret data must reside in L1D
- → Just have a little bit of patience!
- ↗ High impact!



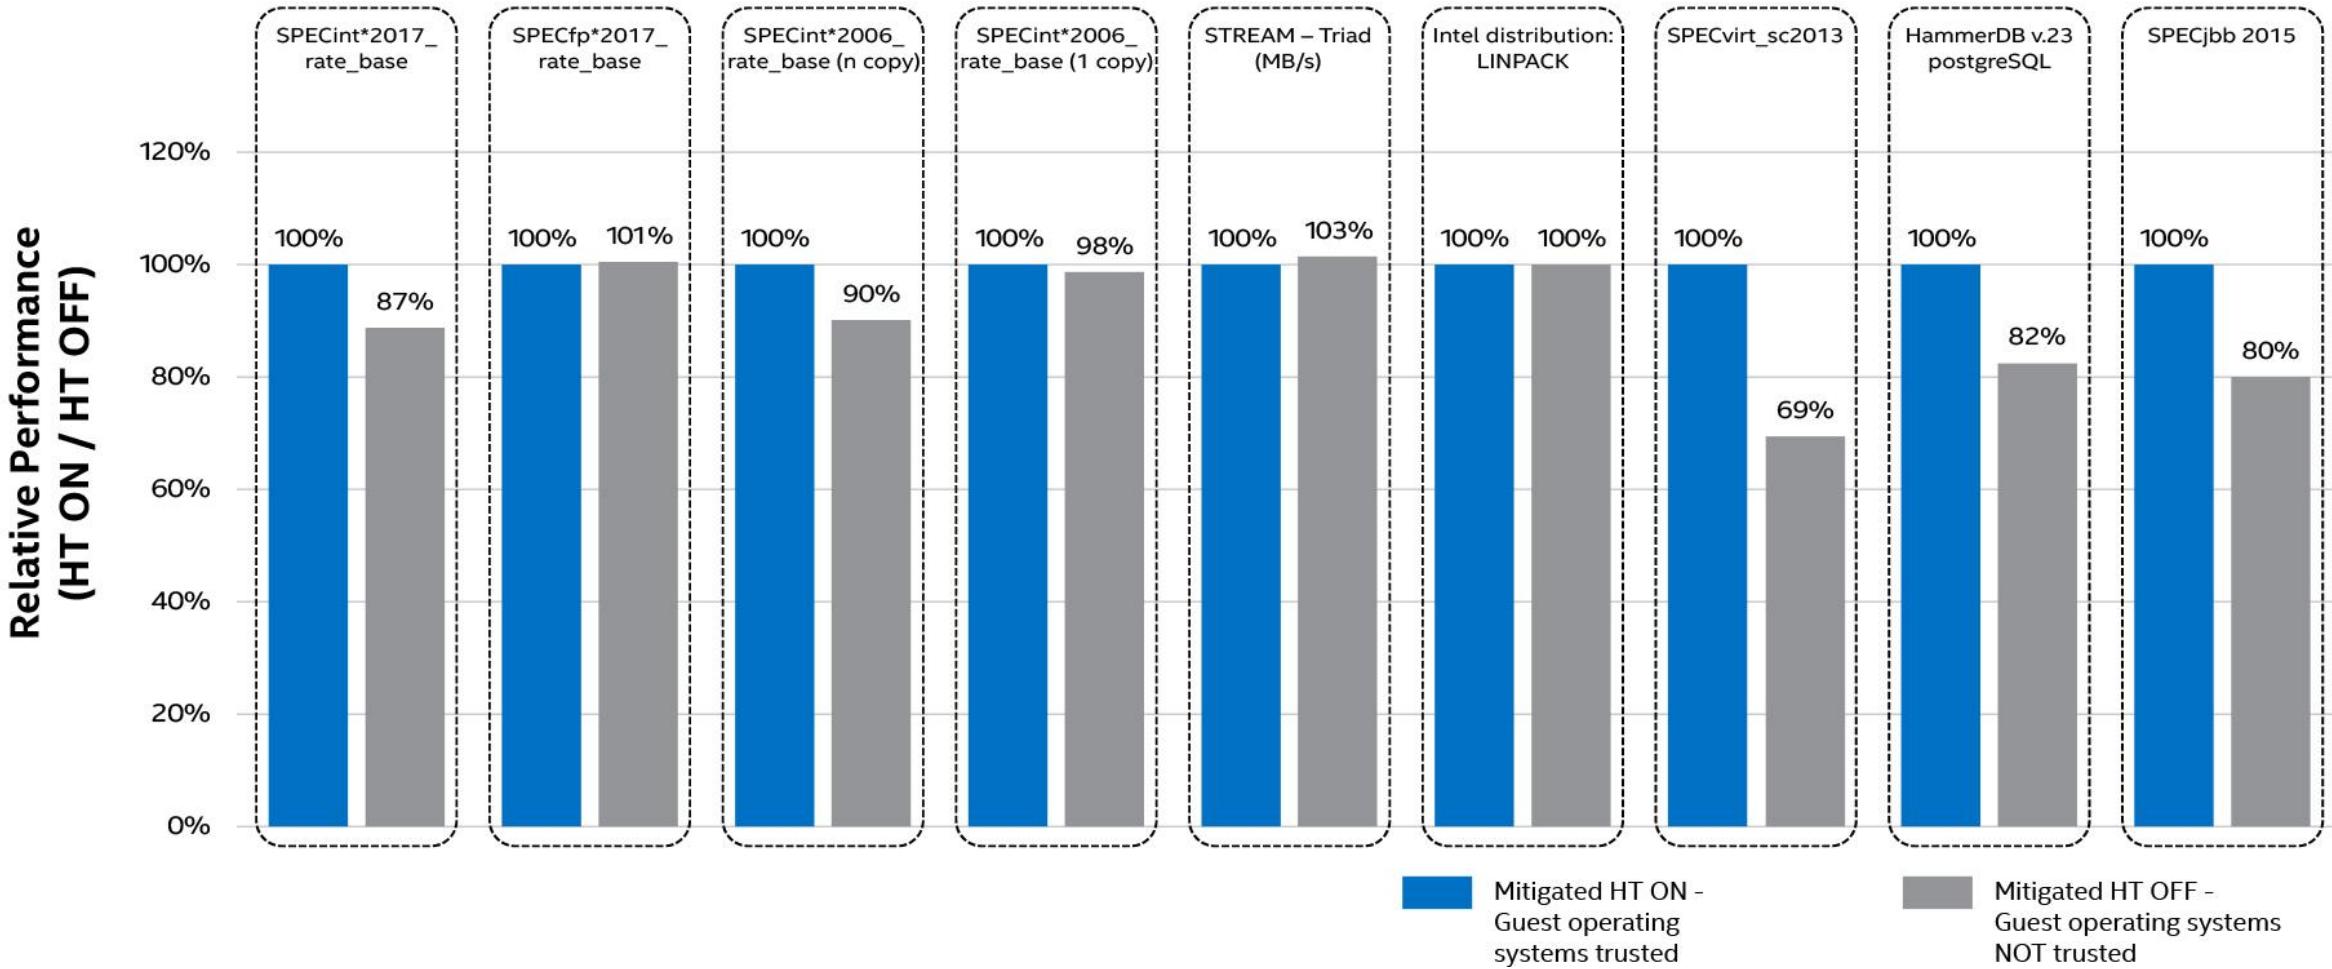
Mitigations

- Long term: Replace chips!
- Short term:
 - Make sure no secrets are in L1D cache
 - → Flush L1D before upon VM-entry
 - → Make sure no two different VMs execute on same physical core
 - Disable HyperThreading
 - Patch VM scheduler



Source: xkcd.com

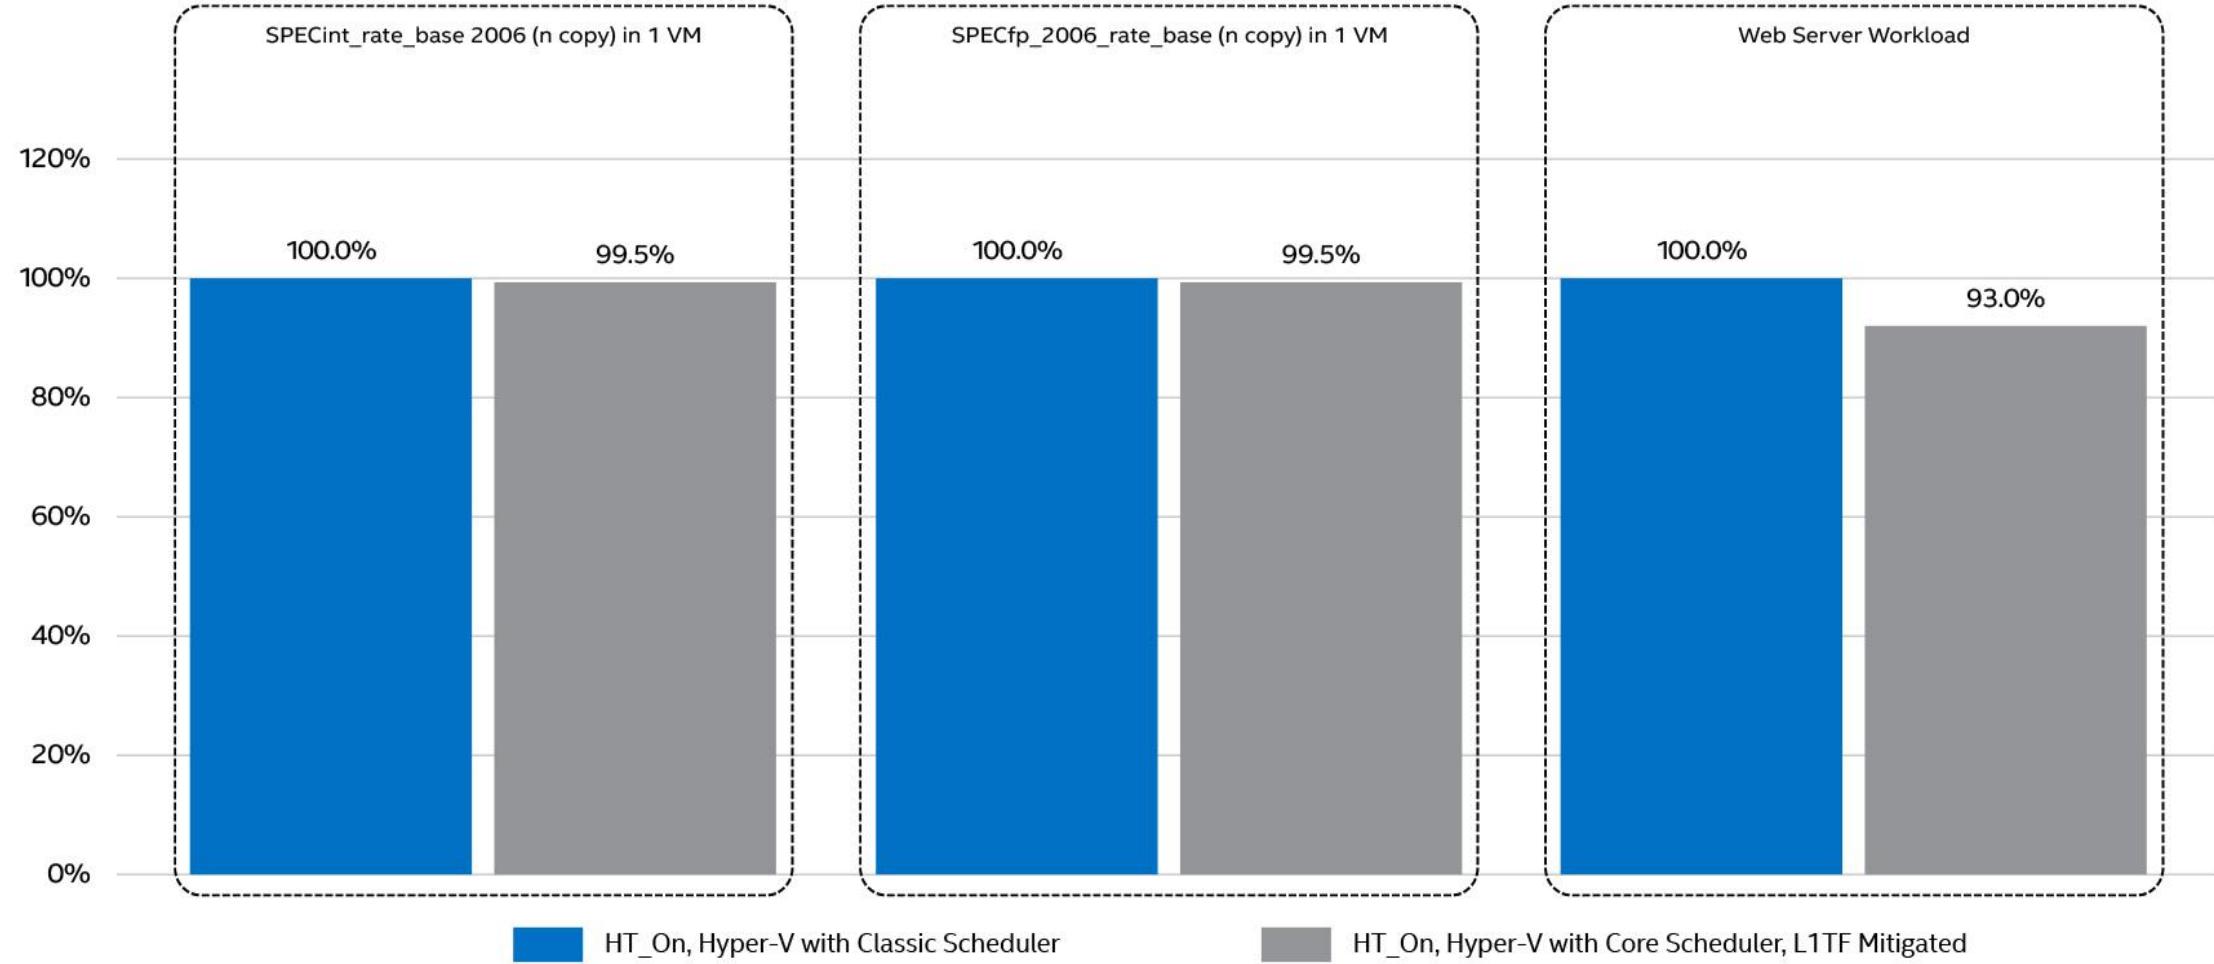
Mitigations – Disabling HyperThreading



source: <https://www.intel.com/content/www/us/en/architecture-and-technology/l1tf.html>

Mitigations – Update VM scheduler

Relative Performance (HT enabled)
(Pre Mitigation / Post Mitigation)



source: <https://www.intel.com/content/www/us/en/architecture-and-technology/l1tf.html>

Foreshadow-SGX: Dismantling Intel SGX's security objectives

Side-Channel Attacks

Attacker



Foreshadow-SGX

Victim



SGX enclave memory

Action: manipulate PT
Carrier: cache
changes

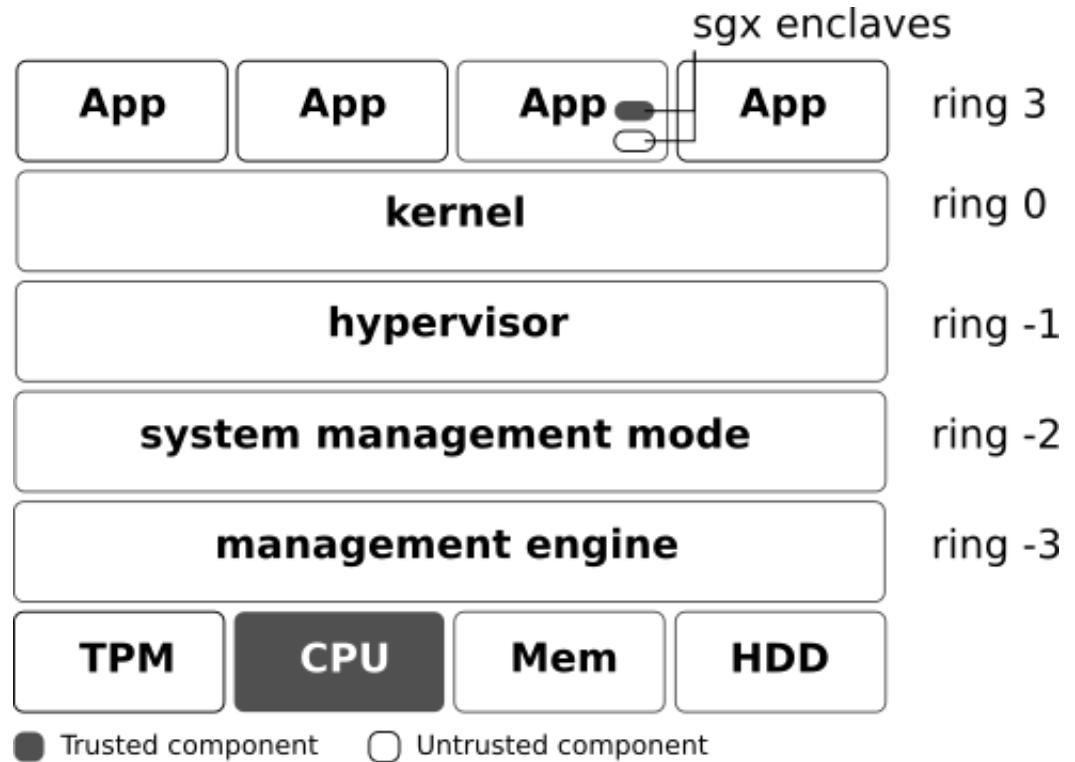
Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

DistrINet

Background: Intel SGX

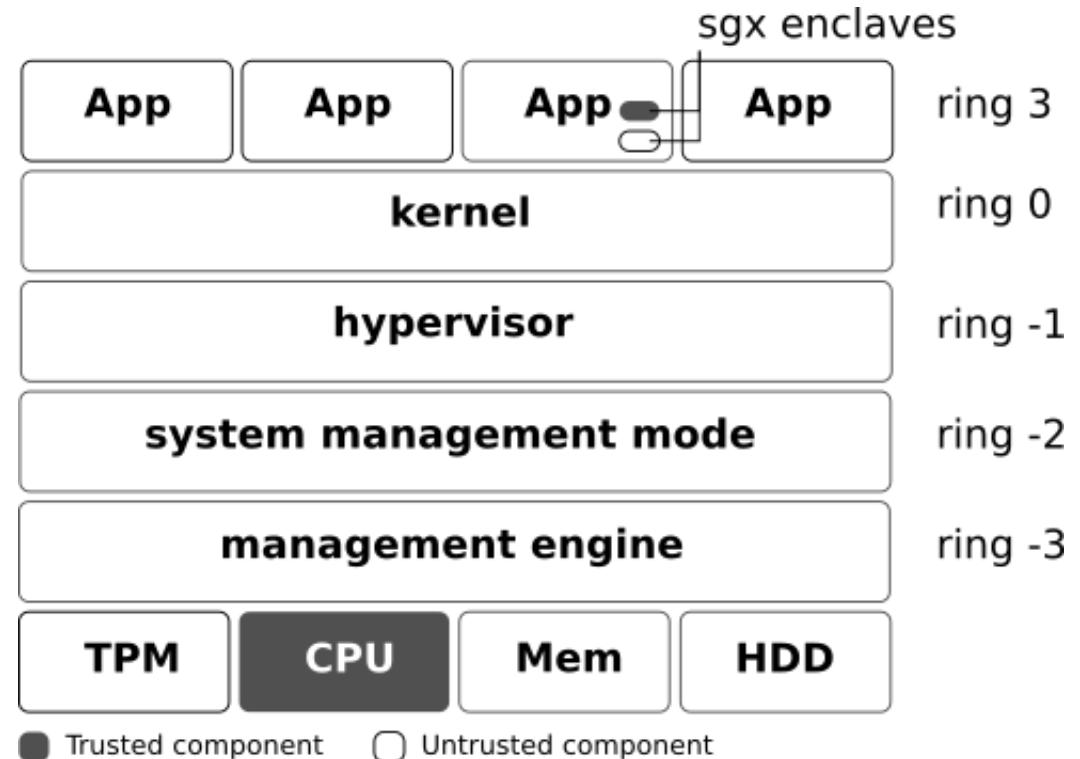
- **Problem:** Huge software TCB
- **Solution:** Protected-Module Architecture (e.g., Intel SGX)
- Only trust Intel hardware/enclaves
- Use cases:
 - protecting finger prints,
 - DRM,
 - password vault,
 - ...



Background: Intel SGX

3 key properties:

- Isolation
- Secure Storage
- Attestation

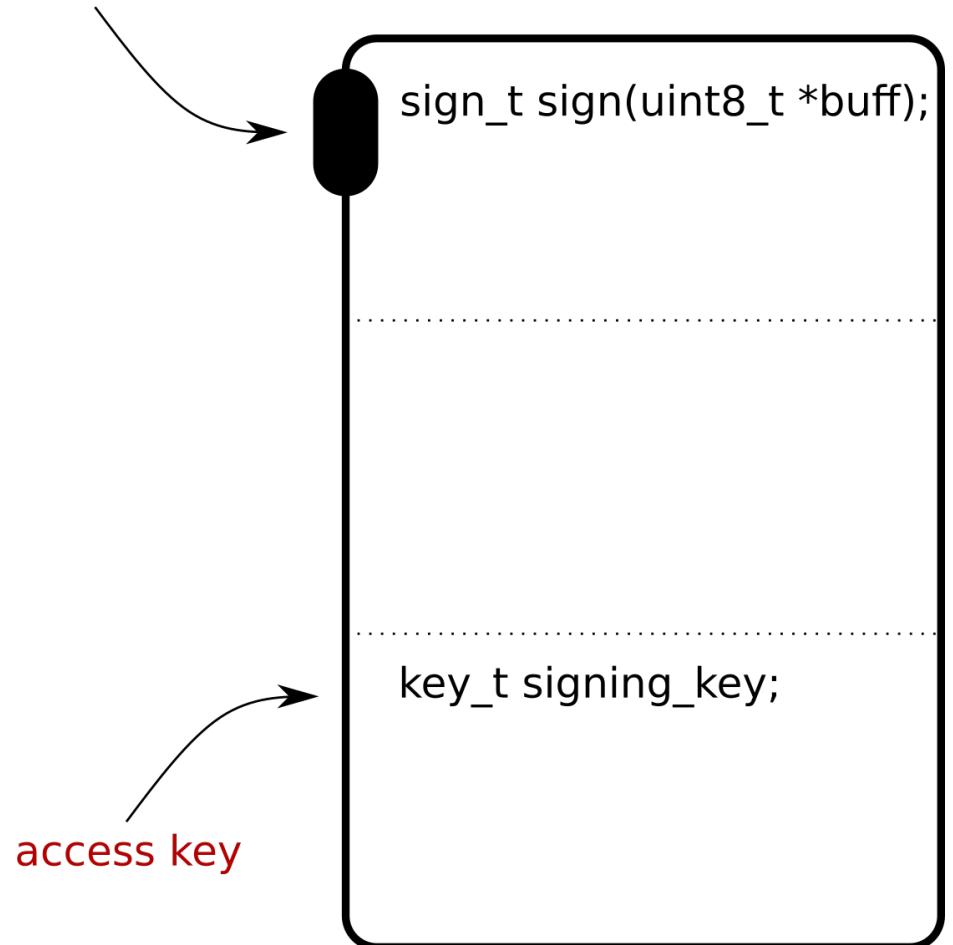


Background: Intel SGX

Isolation

- Enclaves live in process' address space
- Only accessible through specific entry points
- Abort page semantics: Reading enclave memory outside the enclave results in -1.

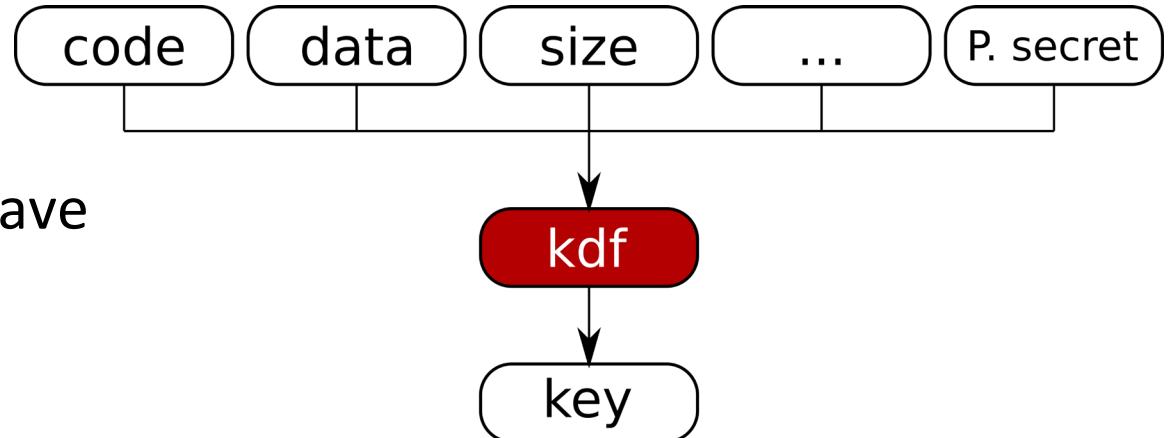
call entrypoint



Background: Intel SGX

Secure Storage:

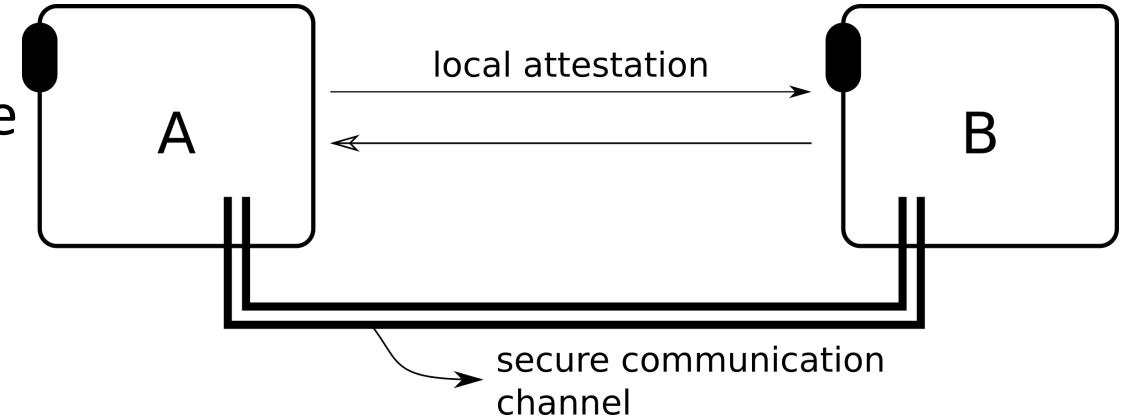
- Enclaves die at power off
- Seal/Unseal confidential data
- Key derivation ensures unique key per enclave
- ↗ Derived keys stored in enclave memory!



Background: Intel SGX

Attestation:

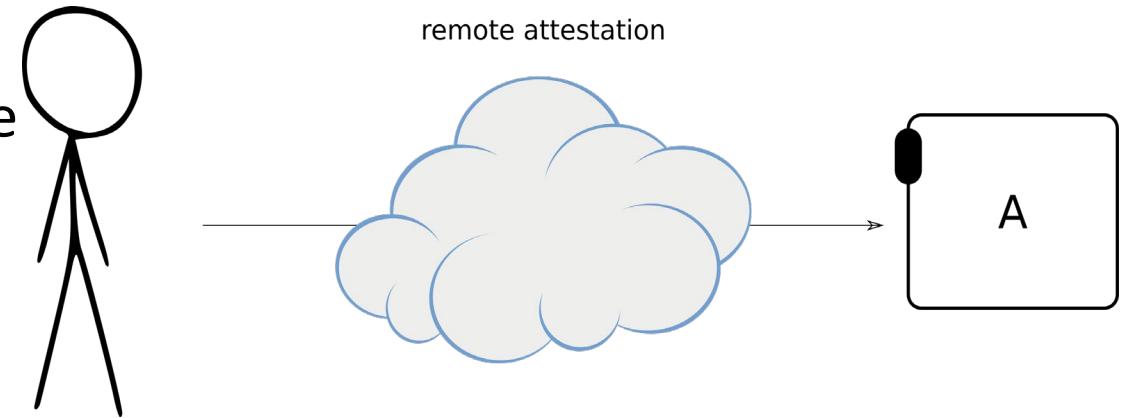
- Prove an enclave has been created correctly
- Both locally as remotely
- Local attestation as a building block for remote attestation



Background: Intel SGX

Attestation:

- Prove an enclave has been created correctly
- Both locally as remotely
- Local attestation as a building block for remote attestation



Side-Channel Attacks

Attacker



Foreshadow-SGX

Victim



SGX enclave memory

Action: manipulate PT
Carrier: cache changes

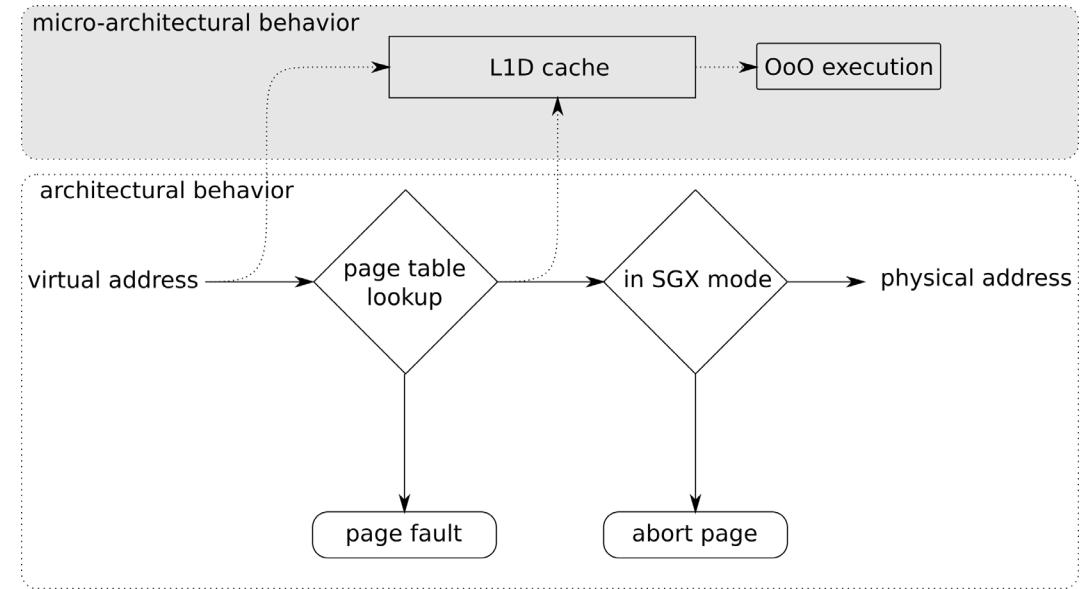
Security flaw: OoO execution leaves traces of transient instructions

Source: xkcd.com

DistrINet

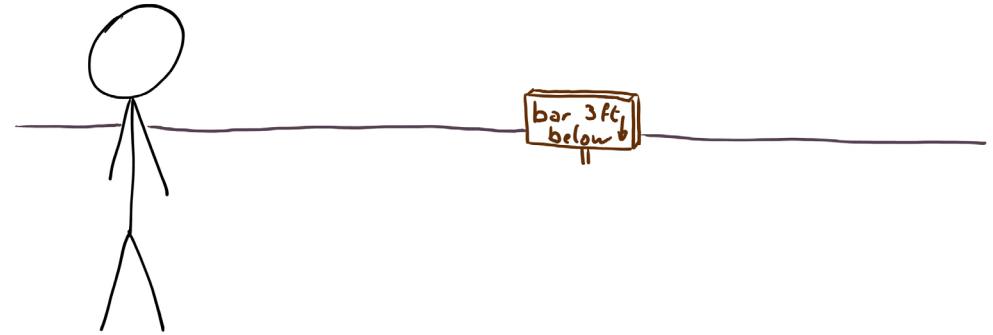
The Attack Approach

- Bypass abort page semantics
- Ensure data in L1D:
 - Single-step through enclave
 - *eldu* instruction as a side effect loads enclave page in L1D!



Impact of Foresight-SGX

- Requirements:
 - Mark enclave page not-present
 - Call enclave/issue “special” SGX instruction
 - Completely breaks enclave isolation, sealed storage and remote/local attestation



Conclusion

- We were lucky Foreshadow was responsibly disclosed
- Foreshadow breaks the virtual memory abstraction
- Modern x86 processors have become too complex to completely understand
- Need to completely rethink past design decisions

Questions?

Raoul Strackx

raoul.strackx@cs.kuleuven.be

[@raoul_strackx](https://twitter.com/raoul_strackx)